

Multivariate Statistical Analysis

Homework 5

Lucas Fellmeth, Helen Kafka, Sven Bergmann

04/04/24

```
library(nnet)
library(caret)
library("e1071")
```

```
set.seed(42)
```

Problem 1

The pendigit dataset contains digitalized handwritten digits.

```
pendigit <- read.csv(file = "../Data_csv/pendigits.csv")
pendigit$digit <- factor(pendigit$digit)
```

The variables $x_1, y_1, \dots, x_8, y_8$ are the coordinates of a pen on a writing pad at eight different time points (so, if you want to visualize the written digit, you have to do `plot(c(x1, ..., x8), c(y1, ..., y8), type='l')`. The variable digit identifies the digit that was written. The goal is to construct a classifier that will identify the handwritten digits as accurately as possible. Split the data into training and test sets (roughly an 80/20 split).

```
i_test_pendigit <- sample(seq_along(pendigit[, 1]), size = length(pendigit[,
1]) * 0.2)
pendigit_test <- pendigit[i_test_pendigit, ]
pendigit_train <- pendigit[-i_test_pendigit, ]
```

Fit single-layer neural networks to the training data, with one, two and three hidden nodes (or more if necessary).

```
num_hidden_nodes_pendigit <- 10
nnets_pendigit <- vector(mode = "list", length = num_hidden_nodes_pendigit)
for (i in 1:num_hidden_nodes_pendigit) {
  nnets_pendigit[[i]] <- nnet(digit ~ . - digit, data = pendigit_train,
    size = i, trace = F)
}
```

Compute the respective misclassification rates on the test set.

```

cms_pendigit <- vector(mode = "list", length = num_hidden_nodes_pendigit)
mscrs_pendigit <- vector(mode = "list", length = num_hidden_nodes_pendigit)
for (i in 1:num_hidden_nodes_pendigit) {
  predictions <- factor(predict(nnets_pendigit[[i]], newdata = pendigit_test,
    type = "class"), levels = levels(pendigit_test$digit))
  cms_pendigit[[i]] <- confusionMatrix(predictions, pendigit_test$digit)
  mscrs_pendigit[[i]] <- 1 - cms_pendigit[[i]]$overall["Accuracy"]
}

```

What's the lowest misclassification rate attained?

```

for (i in 1:num_hidden_nodes_pendigit) {
  cat("Misclassification rate for", i, "hidden nodes:", mscrs_pendigit[[i]],
    "\n")
}

```

```

## Misclassification rate for 1 hidden nodes: 0.9044586
## Misclassification rate for 2 hidden nodes: 0.8080073
## Misclassification rate for 3 hidden nodes: 0.4727025
## Misclassification rate for 4 hidden nodes: 0.333485
## Misclassification rate for 5 hidden nodes: 0.7916288
## Misclassification rate for 6 hidden nodes: 0.4745223
## Misclassification rate for 7 hidden nodes: 0.2961783
## Misclassification rate for 8 hidden nodes: 0.1519563
## Misclassification rate for 9 hidden nodes: 0.1533212
## Misclassification rate for 10 hidden nodes: 0.3416742

```

```

min_mscrs_pendigit <- which.min(mscrs_pendigit)
cat("The lowest misclassification rate is on the neural network \n with",
  min_mscrs_pendigit, "hidden nodes and a rate of", mscrs_pendigit[[min_mscrs_pendigit]],
  ".")

```

```

## The lowest misclassification rate is on the neural network
## with 8 hidden nodes and a rate of 0.1519563 .

```

```

print(cms_pendigit[[min_mscrs_pendigit]])

```

```

## Confusion Matrix and Statistics

```

```

##
##           Reference
## Prediction  0   1   2   3   4   5   6   7   8   9
##           0 200   0   0   0   0   4   0   0   7   2
##           1   1 127   4   1   2   0   0  26   8  11
##           2   0   75 244   0   0   0   0   9   0   0
##           3   2   4   0 193   0   2   0   0   3  19
##           4   0   0   0   1 214   3   1   0   0   0
##           5   0   4   0   0   3 204   2   2  59   4
##           6   4   0   0   0   3   0 198   0   0   0
##           7   0   7   3  10   0   1   3 193  14   0
##           8   3   0   0   0   0   1   0   6 101   0
##           9   0   0   0   0  17   3   0   0   0 190

```

```
##
## Overall Statistics
##
##           Accuracy : 0.848
##           95% CI : (0.8323, 0.8628)
##       No Information Rate : 0.1142
##       P-Value [Acc > NIR] : < 2.2e-16
##
##           Kappa : 0.8309
##
## Mcnemar's Test P-Value : NA
##
## Statistics by Class:
##
##           Class: 0 Class: 1 Class: 2 Class: 3 Class: 4 Class: 5
## Sensitivity      0.95238 0.58525 0.9721 0.94146 0.89540 0.93578
## Specificity      0.99346 0.97325 0.9569 0.98495 0.99745 0.96263
## Pos Pred Value   0.93897 0.70556 0.7439 0.86547 0.97717 0.73381
## Neg Pred Value   0.99496 0.95540 0.9963 0.99392 0.98737 0.99271
## Prevalence       0.09554 0.09873 0.1142 0.09327 0.10874 0.09918
## Detection Rate   0.09099 0.05778 0.1110 0.08781 0.09736 0.09281
## Detection Prevalence 0.09691 0.08189 0.1492 0.10146 0.09964 0.12648
## Balanced Accuracy 0.97292 0.77925 0.9645 0.96321 0.94642 0.94920
##
##           Class: 6 Class: 7 Class: 8 Class: 9
## Sensitivity      0.97059 0.81780 0.52604 0.84071
## Specificity      0.99649 0.98063 0.99501 0.98986
## Pos Pred Value   0.96585 0.83550 0.90991 0.90476
## Neg Pred Value   0.99699 0.97814 0.95640 0.98189
## Prevalence       0.09281 0.10737 0.08735 0.10282
## Detection Rate   0.09008 0.08781 0.04595 0.08644
## Detection Prevalence 0.09327 0.10510 0.05050 0.09554
## Balanced Accuracy 0.98354 0.89921 0.76053 0.91528
```

From the cross-classification tables, which digits have the largest misclassification rates?

```
mscr_digits_best <- vector(mode = "list", length = num_hidden_nodes_pendigit)
for (i in 1:num_hidden_nodes_pendigit) {
  mscr_digits_best[[i]] <- 1 - (cms_pendigit[[min_mscr_pendigit]]$table[i,
    i]/sum(cms_pendigit[[min_mscr_pendigit]]$table[, i]))
}
cat("Misclassification rates for the digits of the set with ",
  min_mscr_pendigit, "\n hidden nodes with the lowest overall rate:\n")
```

```
## Misclassification rates for the digits of the set with 8
## hidden nodes with the lowest overall rate:
```

```
for (i in 1:num_hidden_nodes_pendigit) {
  cat("Misclassification rate of digit ", i - 1, ":", mscr_digits_best[[i]],
    "\n")
}
```

```
## Misclassification rate of digit 0 : 0.04761905
## Misclassification rate of digit 1 : 0.4147465
```

```
## Misclassification rate of digit 2 : 0.02788845
## Misclassification rate of digit 3 : 0.05853659
## Misclassification rate of digit 4 : 0.1046025
## Misclassification rate of digit 5 : 0.06422018
## Misclassification rate of digit 6 : 0.02941176
## Misclassification rate of digit 7 : 0.1822034
## Misclassification rate of digit 8 : 0.4739583
## Misclassification rate of digit 9 : 0.159292
```

```
cat("We have the largest misclassification rate at digit", which.max(mscrs_digits_best) -
    1, ".")
```

```
## We have the largest misclassification rate at digit 8 .
```

Problem 2

The spambase dataset contains data for 4,601 emails which are classified as spam or not spam (as indicated by the variable class);

```
spambase <- read.csv(file = "../Data_csv/spambase.csv")
spambase$class <- factor(spambase$class)
```

58 feature variables are measured on each email. A more detailed description of the data is given on p. 259 of the book. Split the data into training and test sets (roughly an 80/20 split).

```
i_test_spambase <- sample(seq_along(spambase[, 1]), size = length(spambase[,
    1]) * 0.2)
spambase_test <- spambase[i_test_spambase, ]
spambase_train <- spambase[-i_test_spambase, ]
```

Fit a linear support vector machine classifier to the training data, starting with a very large (“infinite”) cost, in the event the groups are separable, and progressively lowering the cost if they aren’t.

We omitted the following code with `cost=Inf` because there was an error message: “Error in svm.default(x, y, scale = scale, ..., na.action = na.action): NA/NaN/Inf in foreign function call (arg 12)”

```
# svmfit_inf_spambase <- svm(class ~ . - class, data =
# spambase_train, cost = Inf, kernel = 'linear')
# summary(svmfit_inf_spambase)
```

```
num_iterations_spambase <- 10
svmfits_spambase <- vector(mode = "list", length = num_iterations_spambase)
# svmfits_spambase[[num_iterations_spambase + 1]] <-
# svmfit_inf_spambase
for (i in num_iterations_spambase:1) {
  svmfits_spambase[[i]] <- svm(class ~ . - class, data = spambase_train,
    cost = 10^i)
}
```

Compute the respective misclassification rates on the test set.

```

cms_spambase <- vector(mode = "list", length = num_iterations_spambase)
mscrs_spambase <- vector(mode = "list", length = num_iterations_spambase)
for (i in 1:num_iterations_spambase) {
  predictions <- factor(predict(svmfits_spambase[[i]], newdata = spambase_test,
    type = "class"), levels = levels(spambase_test$class))
  cms_spambase[[i]] <- confusionMatrix(predictions, spambase_test$class)
  mscrs_spambase[[i]] <- 1 - cms_spambase[[i]]$overall["Accuracy"]
}

```

```

for (i in 1:num_iterations_spambase) {
  cat("Misclassification rate for a svm with the cost", 10^i,
    "is:", mscrs_spambase[[i]], "\n")
}

```

```

## Misclassification rate for a svm with the cost 10 is: 0.009782609
## Misclassification rate for a svm with the cost 100 is: 0.009782609
## Misclassification rate for a svm with the cost 1000 is: 0.009782609
## Misclassification rate for a svm with the cost 10000 is: 0.009782609
## Misclassification rate for a svm with the cost 1e+05 is: 0.009782609
## Misclassification rate for a svm with the cost 1e+06 is: 0.009782609
## Misclassification rate for a svm with the cost 1e+07 is: 0.009782609
## Misclassification rate for a svm with the cost 1e+08 is: 0.009782609
## Misclassification rate for a svm with the cost 1e+09 is: 0.009782609
## Misclassification rate for a svm with the cost 1e+10 is: 0.009782609

```

What's the lowest misclassification rate attained?

```

cat("The lowest misclassificationrate is on the svm \n with a cost of",
  10^which.min(mscrs_spambase), "and a rate of", mscrs_spambase[[which.min(mscrs_spambase)]],
  ".")

```

```

## The lowest misclassificationrate is on the svm
## with a cost of 10 and a rate of 0.009782609 .

```

```

print(cms_spambase[[which.min(mscrs_spambase)]])

```

```

## Confusion Matrix and Statistics
##
##              Reference
## Prediction email spam
##      email    570    7
##      spam      2   341
##
##              Accuracy : 0.9902
##              95% CI : (0.9815, 0.9955)
##      No Information Rate : 0.6217
##      P-Value [Acc > NIR] : <2e-16
##
##              Kappa : 0.9791
##
##      McNemar's Test P-Value : 0.1824

```

```
##
##      Sensitivity : 0.9965
##      Specificity : 0.9799
##      Pos Pred Value : 0.9879
##      Neg Pred Value : 0.9942
##      Prevalence : 0.6217
##      Detection Rate : 0.6196
##      Detection Prevalence : 0.6272
##      Balanced Accuracy : 0.9882
##
##      'Positive' Class : email
##
```